Predicting average ratings for board games

A case study in futility

The name of the game

- BoardGameGeek.com
- Created in 2000
- Over 80,000 board games
- Strong member community

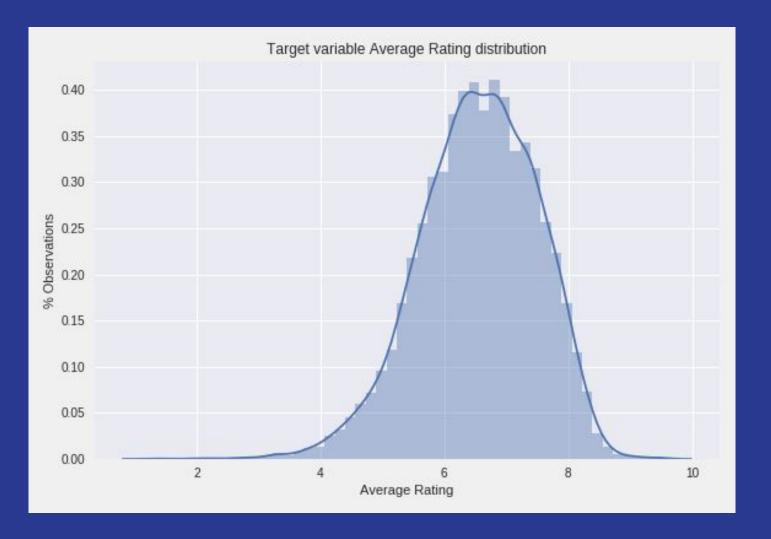


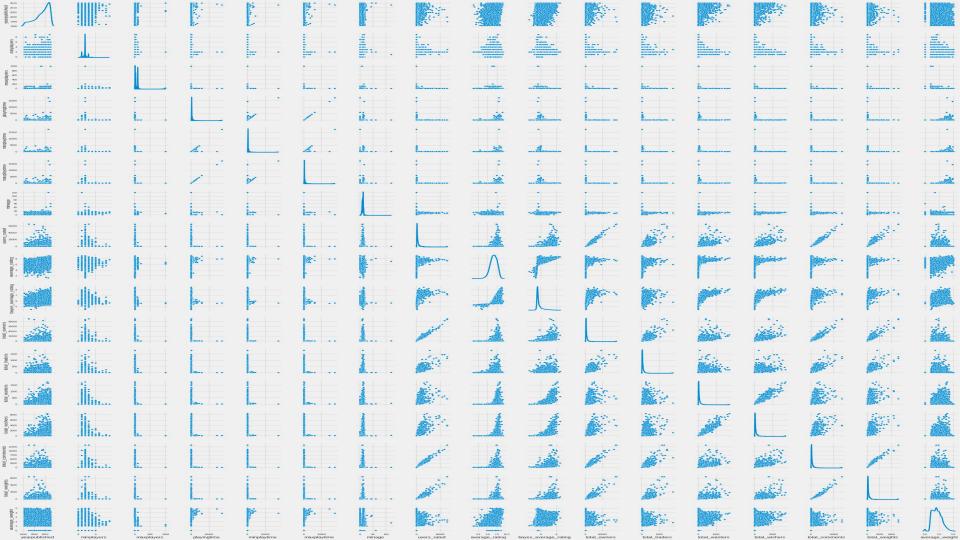


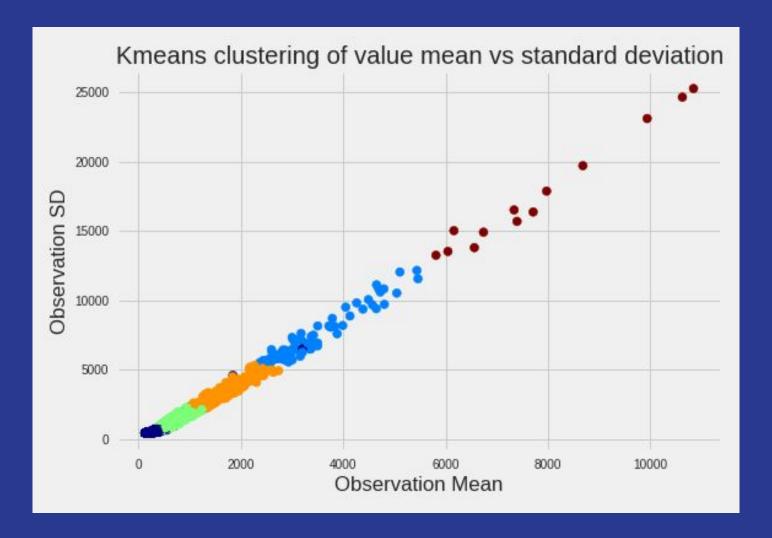
1. How to play

I used Scrapy and Beautiful Soup to pull down my data into csv form before pulling it into pandas.

- → Cleaning
 "Empty" games
- → Simplify
 Non-numerical features
- → Explore...









2. Correlations

yearpublished	0.348639
minplayers	-0.169357
maxplayers	-0.024367
playingtime	0.118681
minplaytime	0.114625
maxplaytime	0.118681
minage	0.141881
users_rated	0.153301
average rating	1.000000
bayes average rating	0.516445
total owners	0.206832
total_traders	0.081275
total wanters	0.318231
total_wishers	0.257790
total comments	0.152957
total_weights	0.133433
average_weight	0.510265
Name: average_rating,	dtype: float64



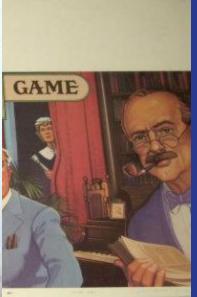


3. Regression

I removed all collinear features before modeling my training data.

- ightharpoonup Cross-valididated L1 100 sample α
- → Cross-validated L2
- Optimized CV L1
 500 sample α
- Lasso
 Trained on the entire sample using the optimized α









3. Scoring

- → MSE
 - 0.459
- → RMSE
 - 0.679
- \rightarrow Target σ
 - 0.953



